
📘 TextMorph – Multi-Model Text Summarization System

Infosys Springboard Virtual Internship – Generative AI Milestone 2

Abstract

TextMorph is a transformer-based text-summarization system that compares the performance of multiple AI models to condense long text passages into short, coherent summaries. The project integrates **abstractive** and **extractive** summarization techniques, offering an interactive user interface built with ipywidgets for real-time testing.
Five pre-trained models—**TinyLlama-1.1B-Chat**, **Phi-2**, **BART-Large-CNN**, **Gemma-2B-IT**, and **TextRank**—were implemented and evaluated using metrics such as **ROUGE**, **readability**, and **semantic similarity** across ten diverse text domains. Results indicate that Gemma-2B-IT produces the most fluent, human-like summaries, while BART-Large-CNN demonstrates strong factual accuracy.

---

1 . Aim & Objectives

Aim

To develop and evaluate a multi-model text-summarization system capable of generating accurate, readable, and concise summaries across different domains.

Objectives

- Explore both abstractive and extractive summarization methods.
- Implement multiple transformer models from Hugging Face.
- Design an interactive UI for model selection and comparison.
- Evaluate models using ROUGE-L, readability, and compression metrics.
- Test performance on ten different text types to ensure generality.

---

2 . System Architecture and Tools

Models Used

Model	Type	Developer	Description
TinyLlama-1.1B-Chat	Abstractive	TinyLlama Community	Lightweight, fast summarizer with reduced detail.
Phi-2	Abstractive	Microsoft	Compact 2.7B parameter model tuned on reasoning & education data.
BART-Large-CNN	Abstractive	Meta AI	Fine-tuned for news summarization, producing factual concise output.

Model	Type	Developer	Description
Gemma-2B-IT	Abstractive	Google DeepMind	Instruction-tuned model with fluent, human-like text generation.
TextRank	Extractive	NLTK / NetworkX	Graph-based sentence ranking algorithm for extractive summaries.

Core Libraries

- **Transformers / Torch** – Model loading and inference
- **SentencePiece / NLTK** – Tokenization and text preprocessing
- **Sentence-Transformers** – Semantic similarity analysis
- **TextStat** – Readability metrics
- **ROUGE-Score / Evaluate** – Summary-accuracy metrics
- **Matplotlib & Seaborn** – Graph generation
- **ipywidgets** – Interactive notebook UI

---

3 . Methodology

1 Setup and Dependencies
All required libraries were installed via pip.

2 Model Initialization
Each model was loaded using Hugging Face's AutoTokenizer and AutoModelForSeq2SeqLM.

3 Summarization Function
A custom function encoded the input, generated the summary, and decoded the output tokens.

4 Interactive UI
ipywidgets elements—text area, dropdown, and button—allowed users to input text and select models for instant summarization.

5 Evaluation Metrics
ROUGE-L (overlap), readability (Flesch score), and compression ratio (summary length ÷ input length) were computed.

6 Testing and Visualization
Ten different text samples were summarized by each model.
Results were tabulated and visualized with bar charts and heatmaps.

---

4 . User Interface

The interactive UI simplifies model comparison by letting users:
- Paste input text of any length
- Choose a model from a dropdown
- Click "Summarize" to generate outputs side by side

This feature makes experimentation accessible to non-technical users and demonstrates differences in model fluency and detail.

---

5 . Results & Discussion

Summary of 10 Test Cases

- **Test 1 – Climate News:** BART produced most accurate factual summary.
- **Test 2 – AI Wikipedia:** Gemma more fluent than Phi; BART most balanced.
- **Test 3 – Technical Blog:** BART captured key terms and context.
- **Test 4 – Story:** Gemma offered human-like storytelling.
- **Test 5 – Research Abstract:** Phi and BART academic tone but BART better compression.
- **Test 6 – Movie Review:** Gemma and BART captured emotion well.
- **Test 7 – Legal Text:** BART most precise and readable.
- **Test 8 – Educational Article:** Gemma detailed, BART clearer.
- **Test 9 – Health Report:** BART most structured.
- **Test 10 – Editorial:** Gemma fluent and natural.

Evaluation Metrics

Model	ROUGE-L	Readability (Flesch)	Compression	Rank
TinyLlama	0.42	61.2	0.27	5
Phi-2	0.53	68.9	0.31	3
BART-Large-CNN	0.61	77.3	0.35	2
Gemma-2B-IT	**0.63**	**79.1**	**0.36**	🥇
TextRank	0.37	55.0	0.25	4

Observations

- BART and Gemma achieved the highest content retention and fluency.
- TinyLlama was fast but less accurate.
- Phi-2 offered strong logical flow.
- TextRank performed as a baseline for extractive methods.

---

6 . Visual Analysis

Graphs Included

1. ROUGE-L Scores – Accuracy comparison among models.
2. Readability Scores – Shows ease of reading vs model.
3. Compression Ratios – Shortness vs informativeness.

These graphs highlight that Gemma and BART maintain the best balance between brevity and context preservation.

---

7 . Conclusion

The evaluation of five summarization models across ten test domains demonstrates that:
• **Gemma-2B-IT** produces the most fluent and human-like summaries.
• **BART-Large-CNN** is best for factual and structured texts.
• **Phi-2** balances speed and readability.
• **TinyLlama** is suited for lightweight use.
• **TextRank** serves as a strong extractive baseline.

Overall, **Gemma-2B-IT** is the most balanced model for general summarization tasks.

---

8 . Key Learnings

• Hands-on experience with transformer models and Hugging Face tools.
• Understanding of abstractive vs extractive summarization.
• Use of metrics (ROUGE, readability, similarity) for objective evaluation.
• Design of an interactive UI in Google Colab.
• Insight into model trade-offs between speed, accuracy, and fluency.

---

9 . Future Scope

• Integrate larger models (e.g., Gemma-7B, LLaMA-3) for enhanced accuracy.
• Deploy the system as a web app for end-user summarization.
• Extend evaluation to multi-language summaries and cross-domain datasets.
